

APEX DATA ORGANIZATION AND PROCESSING




CONFIDENTIAL AND PROPRIETARY



SECURE CONNECTIONS
FOR A SMARTER WORLD

Let's start with an image. The first step is to split the image into slices. In this example, we'll use a slice of size 32x4. To process the slice, we must divide it into blocks. We will call these blocks "chunks" hereafter. For each block, represented as multidimensional array, stride measures the number of locations in memory from one row of pixels to the next row of pixels. If padding is presented, the stride will be wider than the width of the core image block.



When splitting the slice into chunks, we need to end up with one chunk for each CU. In this example we'll be using 4 CUs. The width of each chunk is the width of the slice divided by the number of CUs, so here the chunk width will be $32/4 = 8$. The chunk height is the same as the slice height. Therefore our chunk size will be 8x4.

The numbers in each box are the data values. Each red box represents a chunk.

00	01	02	03	04	05	06	07	08	09	0A	0B	0C	0D	0E	0F	10	11	12	13	14	15	16	17	18	19	1A	1B	1C	1D	1E	1F
20	21	22	23	24	25	26	27	28	29	2A	2B	2C	2D	2E	2F	30	31	32	33	34	35	36	37	38	39	3A	3B	3C	3D	3E	3F
40	41	42	43	44	45	46	47	48	49	4A	4B	4C	4D	4E	4F	50	51	52	53	54	55	56	57	58	59	5A	5B	5C	5D	5E	5F
60	61	62	63	64	65	66	67	68	69	6A	6B	6C	6D	6E	6F	70	71	72	73	74	75	76	77	78	79	7A	7B	7C	7D	7E	7F

The data for each chunk will be transferred into the memory of its corresponding CU.

00	01	02	03	04	05	06	07	08	09	0A	0B	0C	0D	0E	0F	10	11	12	13	14	15	16	17	18	19	1A	1B	1C	1D	1E	1F
20	21	22	23	24	25	26	27	28	29	2A	2B	2C	2D	2E	2F	30	31	32	33	34	35	36	37	38	39	3A	3B	3C	3D	3E	3F
40	41	42	43	44	45	46	47	48	49	4A	4B	4C	4D	4E	4F	50	51	52	53	54	55	56	57	58	59	5A	5B	5C	5D	5E	5F
60	61	62	63	64	65	66	67	68	69	6A	6B	6C	6D	6E	6F	70	71	72	73	74	75	76	77	78	79	7A	7B	7C	7D	7E	7F

↓
CU₀

↓
CU₁

↓
CU₂

↓
CU₃


The numbers in blue represent the CMEM addresses. In this example, block stride is 8.

	CU ₀	CU ₁	CU ₂	CU ₃
A+00:	00	08	10	18
A+01:	01	09	11	19
A+02:	02	0A	12	1A
A+03:	03	0B	13	1B
A+04:	04	0C	14	1C
A+05:	05	0D	15	1D
A+06:	06	0E	16	1E
A+07:	07	0F	17	1F
A+08:	20	28	30	38
A+09:	21	29	31	39
A+22:	46	4E	56	5E
A+23:	47	4F	57	5F
A+24:	60	68	70	78
A+25:	61	69	71	79
A+26:	62	6A	72	7A
A+27:	63	6B	73	7B
A+28:	64	6C	74	7C
A+29:	65	6D	75	7D
A+30:	66	6E	76	7E
A+31:	67	6F	77	7F

In image processing, some algorithms have spatial dependencies. This means that when they're processing a section of the image, they need to have access to the data that surrounds it.


Ordinarily, a CU will only contain the data from the chunk that it's currently processing. However if we're dealing with an algorithm that has spatial dependencies, we use padding to ensure that the chunk has access to its surrounding data.

In the figure below, we can see how a spatial dependency of 1 pixel on each side maps out.



- Pixel being processed
- Spatial dependencies


Spatial dependencies add some requirements to properly process the image. The pixels that are on the edge of the image can't be processed because there is no data beyond the border. To solve this problem, we use image border padding that adds padding to the edge of the original image. The data beyond the edge might not exist, so it is often replicated from the edge values while the inner rectangle consists of valid region pixels.



With spatial dependencies, we encounter another problem when dealing with chunks: the kernel needs to have access to data from the neighboring chunks to process the pixels at the edge of the chunks. To do so, we perform chunk padding. Chunk padding reduces the complexity of kernel implementation and increases its reusability.

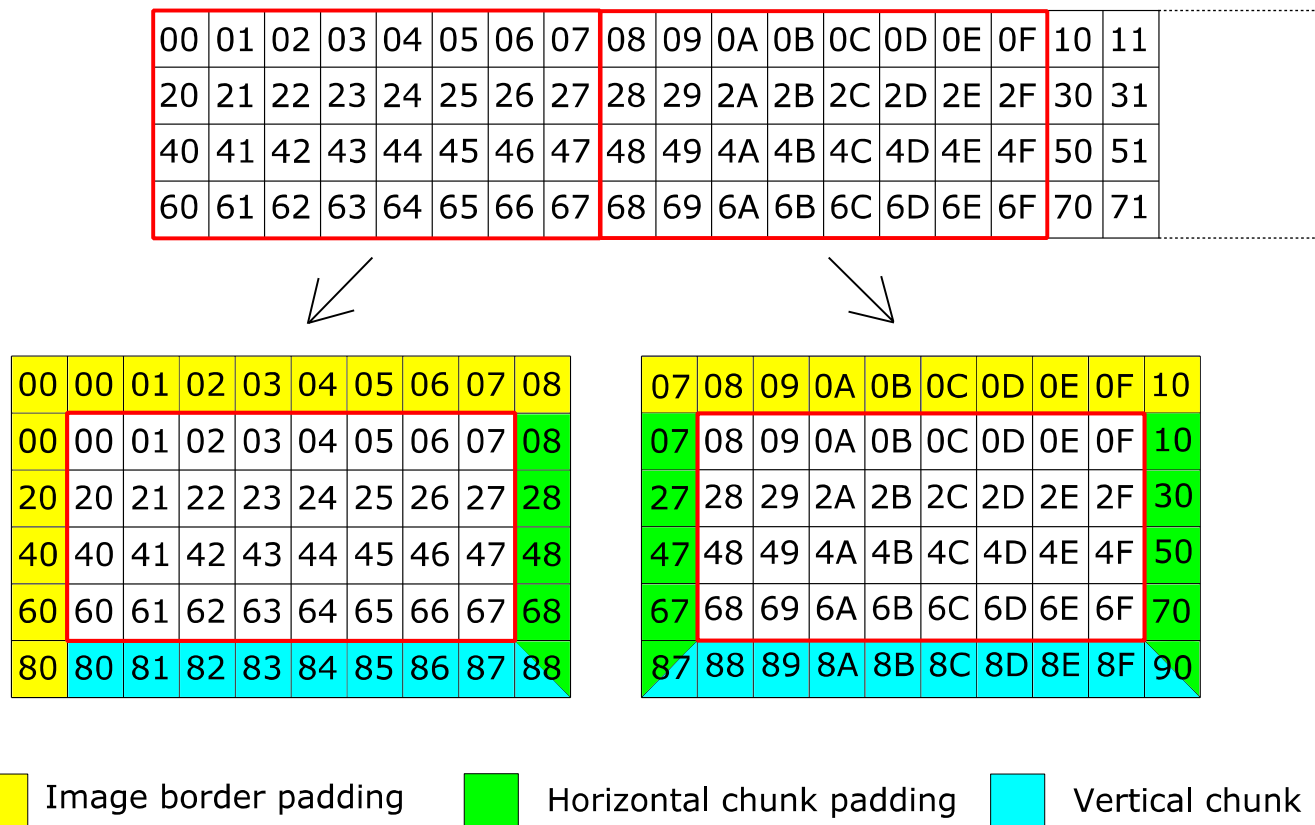
There are two types of chunk padding: vertical padding and horizontal padding.

Vertical padding is used to access data that is above or below the chunk being processed. Since the chunk has a set height, that data will come from the chunk in neighboring slice that is either above or below the current slice.

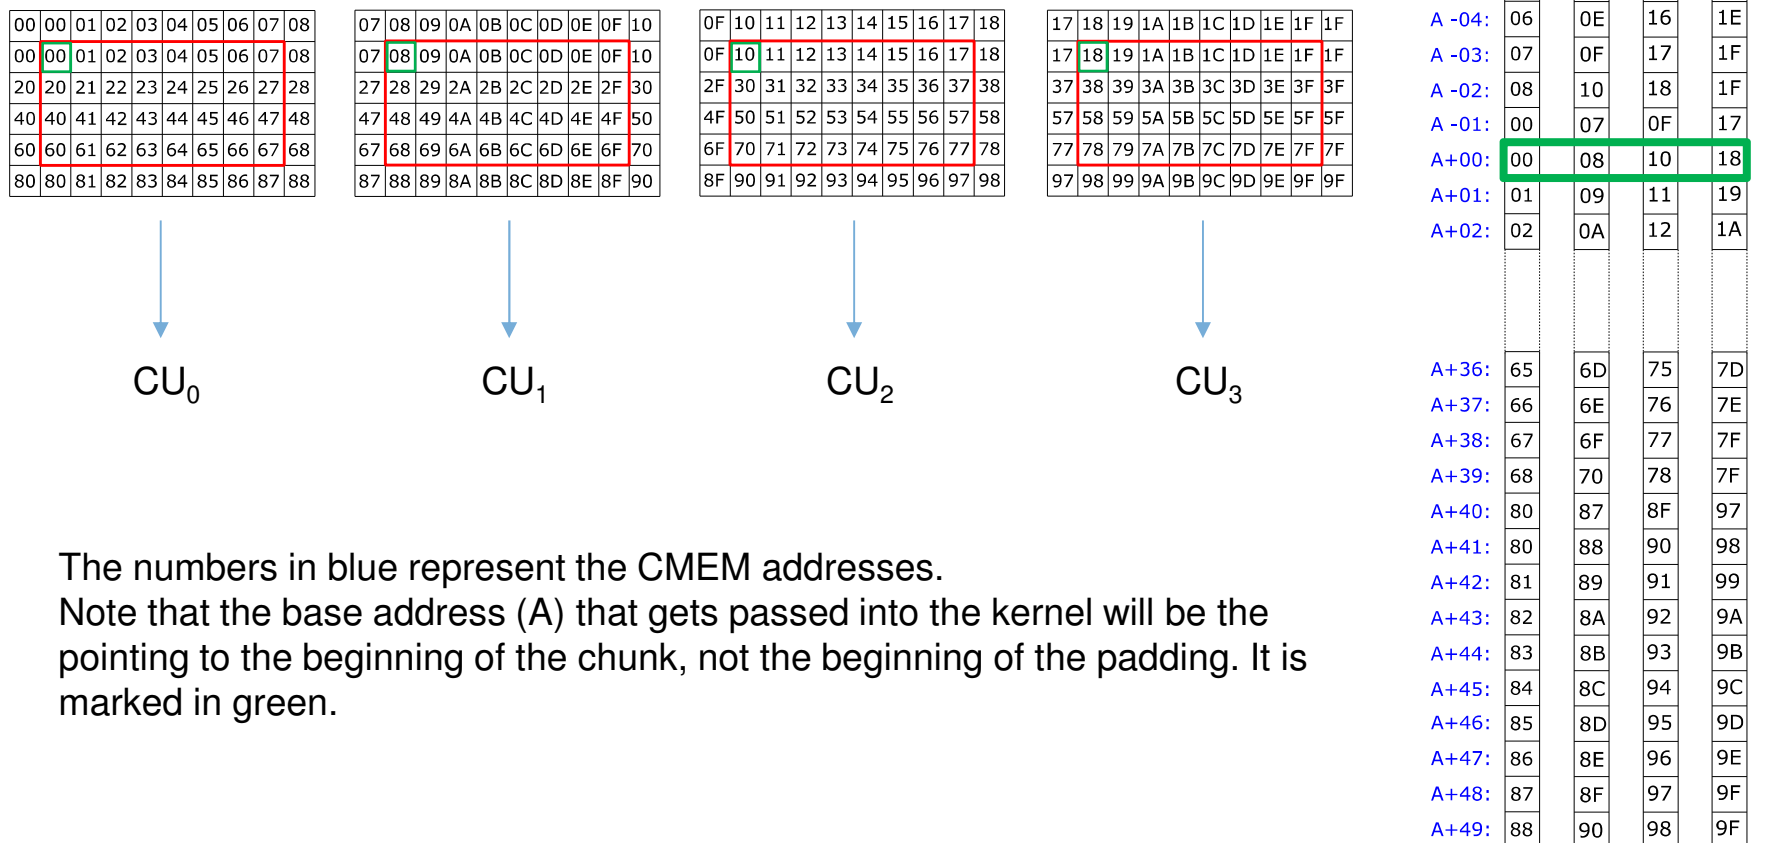


Horizontal padding is used to access data that is to the left or right of the chunk being processed.

Let's go back to our slice example. This time we're going to assume that we have a spatial dependency of 1 pixel on each side. Since this is the topmost slice, we need to pad perform image border padding by pixel replication. We also need to perform chunk padding. Here is what it would look like:



After the transfer, each CU contains the data from its corresponding chunk.
Note that chunk padding is automatically managed and provided by the APEX Core Framework.



The numbers in blue represent the CMEM addresses.
Note that the base address (A) that gets passed into the kernel will be the pointing to the beginning of the chunk, not the beginning of the padding. It is marked in green.





SECURE CONNECTIONS
FOR A SMARTER WORLD